Introduction Big Data

Agenda

- What is Big Data?
- How Big is our Data Universe?
- Sources of Big Data?
- Four V's of Big Data
- Data Storage & Analysis of Big Data
- Hadoop as Solution

What is Big Data?

- **Large Volume of Data both Structured & Unstructured**

Big Data

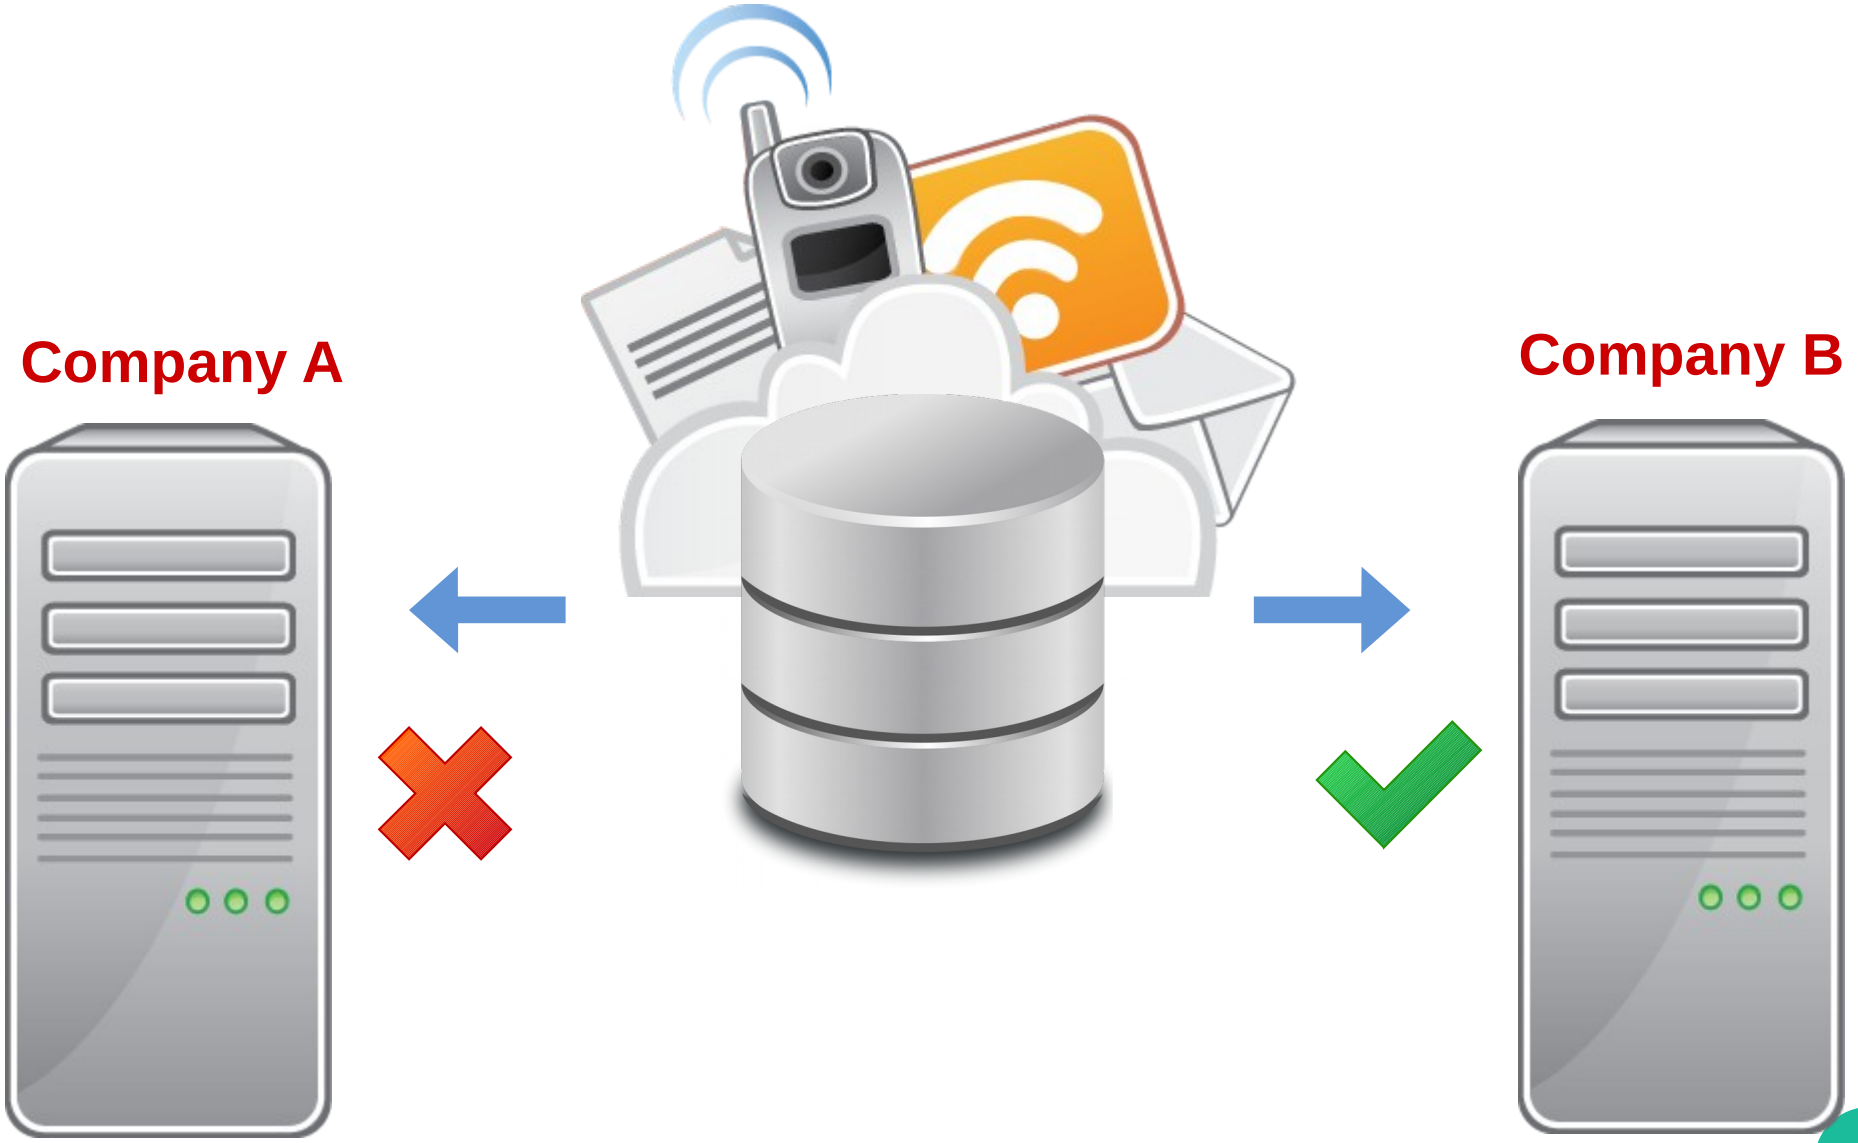

How Big is Our Data Universe?

- Every Day, we create 2.5 quintillion bytes of Data.
- 90% of the data in the world today has been created int the last two years alone.
- The production of data is expanding by 4300% increase in annual data generation by 2020.

How Big is Our Data Universe?

- **IDC estimate “digital universe” at 4.4 zettabytes in 2013 and is forecasting a tenfold growth by 2020 to 44 zettabytes.**

What is Zettabyte?

1 Bit	Binary Digit
8 Bits	1 Byte
1024 Bytes	1 Kilobyte
1024 Kilobytes	1 Megabyte
1024 Megabytes	1 Gigabyte
1024 Gigabytes	1 Terabyte
1024 Terabytes	1 Petabyte
1024 Petabytes	1 Exabyte
1024 Exabytes	1 Zettabyte

Big Data Sources

- The New York Stock Exchange generates about 4−5 terabytes of data per day.
- Facebook hosts more than 240 billion photos, growing at 7 petabytes per month.
- Ancestry.com, the genealogy site, stores around 10 petabytes of data.
- The Internet Archive stores around 18.5 petabytes of data.
- The Large Hadron Collider near Geneva, Switzerland, produces about 30 petabytes of data per year.

Four V's of Big Data

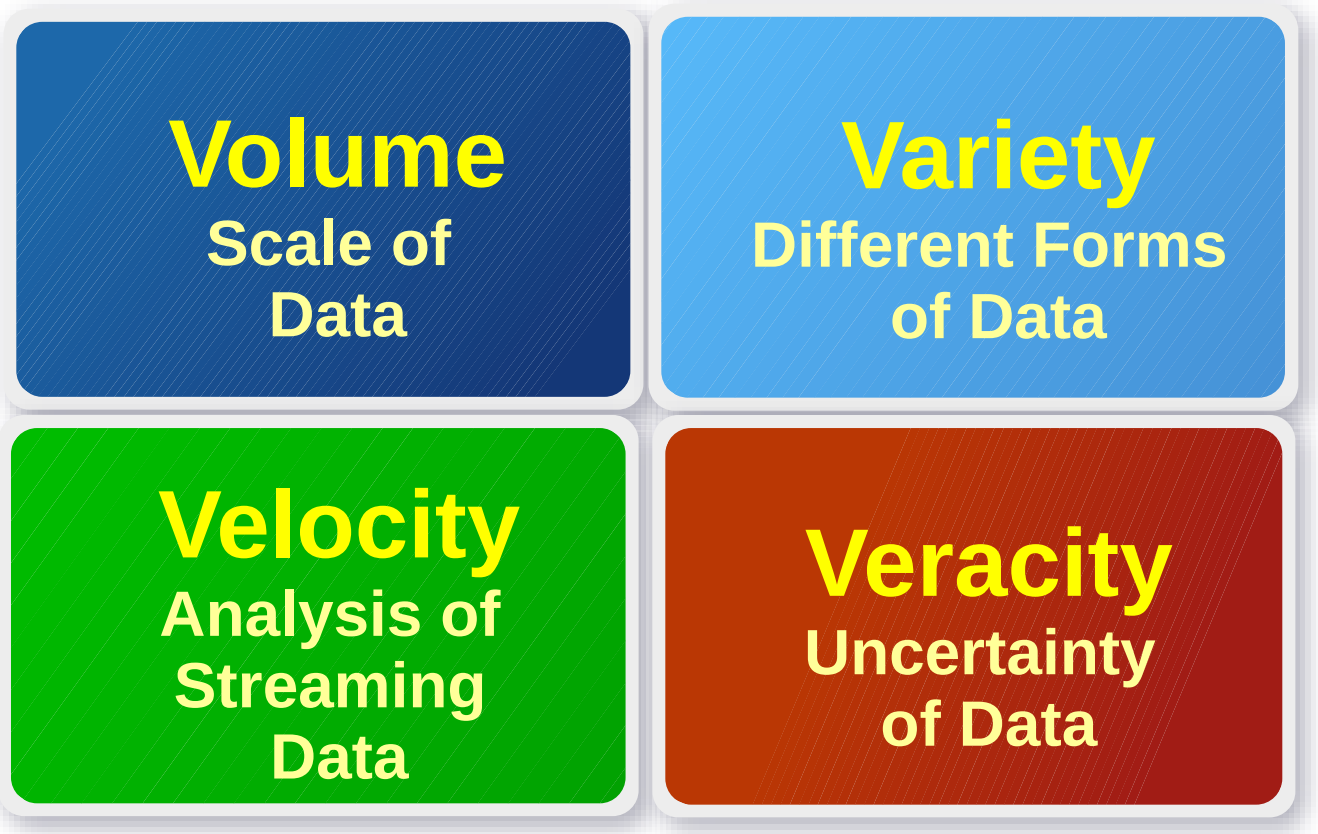

Big Data Problem

Data Storage & Analysis of Big Data

- **Storage capacity has grown exponentially but read speed has not kept up**
  - **1990:**
    - **Store 1,400 MB**
    - **Transfer speed of 4.5MB/s**
    - **Read the entire drive in ~ 5 minutes**
  - **2010:**
    - **Store 1 TB**
    - **Transfer speed of 100MB/s**
    - **Read the entire drive in ~ 3 hours**

Data Storage & Analysis of Big Data

- **Why not reading from multiple disks at once to reduce time?**
  - Imagine 100 drives, each holding one hundredth of the data. Working in parallel
  - Read entire data in under two minutes.
- **Yes, Hadoop does that.**

Problems with Reading from Multiple Disks Together

- **Hardware failure**
  - **HDFS solves this.**
- **Combining data from Multiple Disks after reads**
  - **Map Reduce Solves this.**

Hadoop as Affordable Solution

- **Hadoop provides: a reliable, scalable platform for storage and analysis.**
- **It runs on commodity hardware.**
- **It is open source.**

Hadoop Eco System

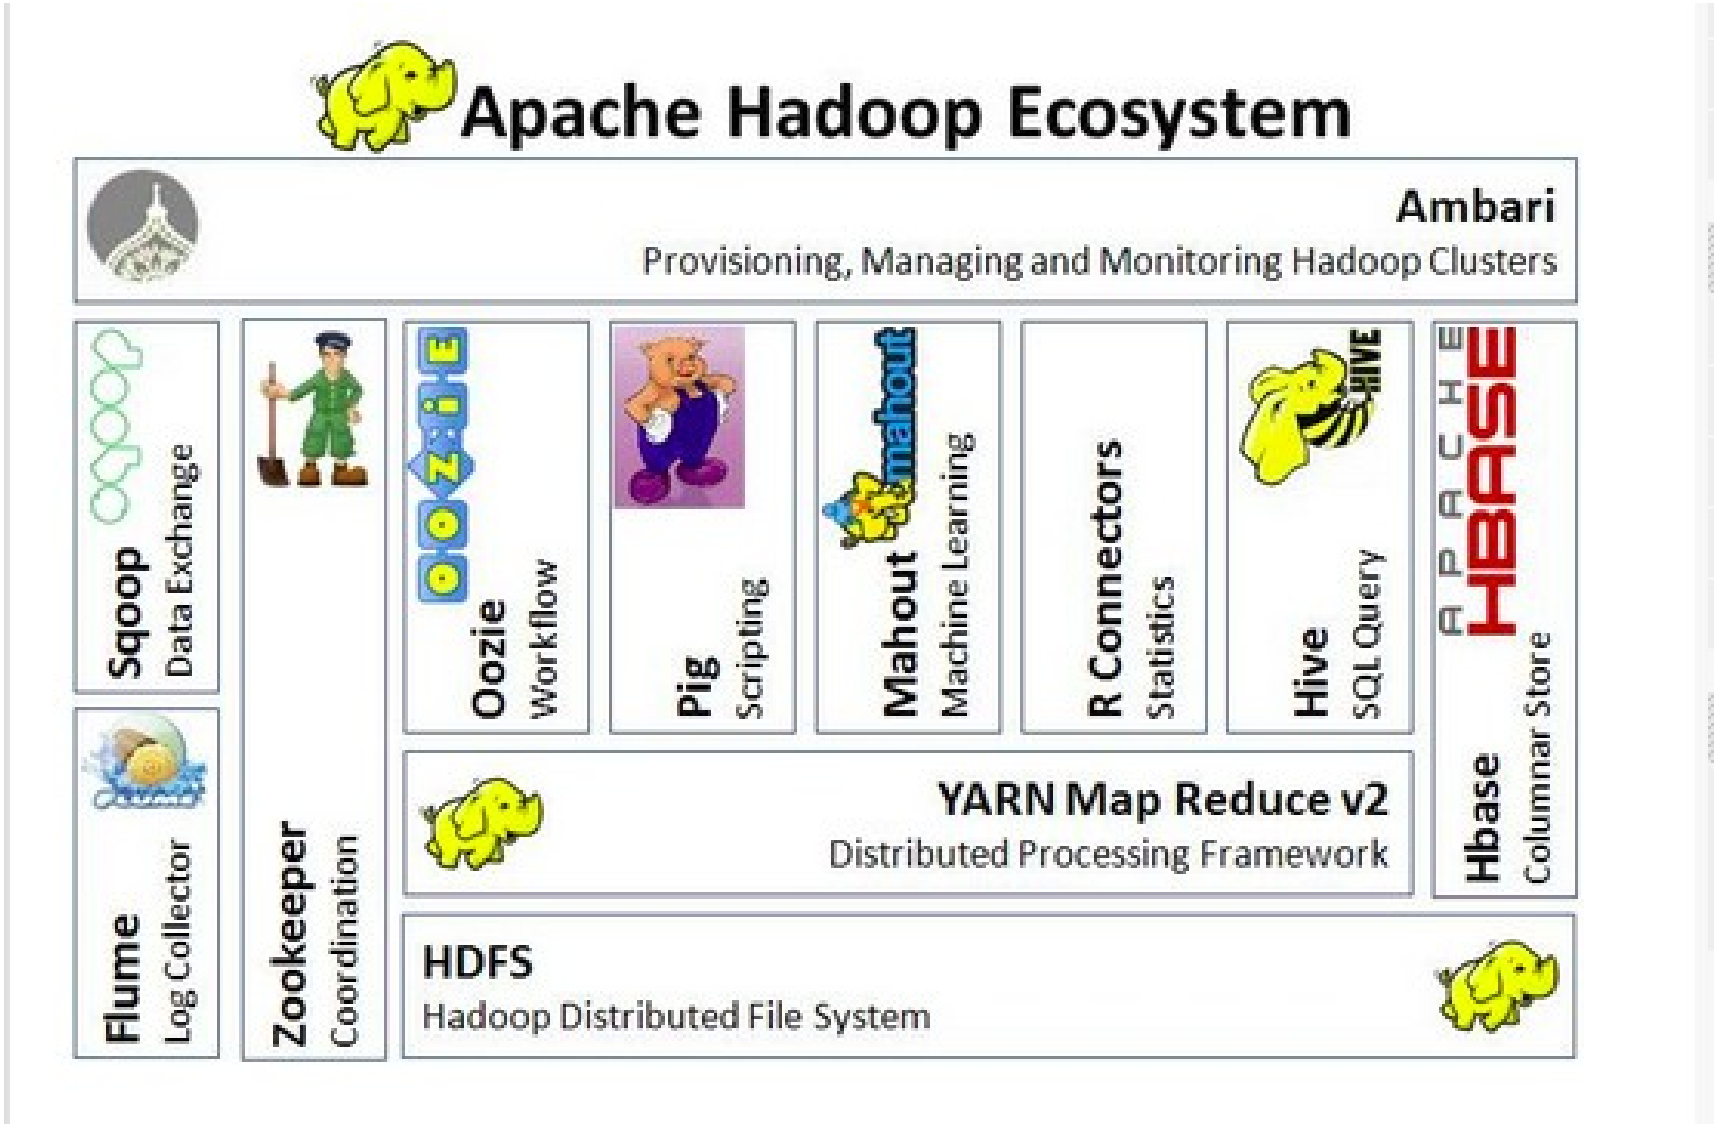

Resources

- **Hadoop: The Definitive Guide**
  - Tom White (Author)
  - O'Reilly Media; 4th Edition.