

HDFS- Deep Dive



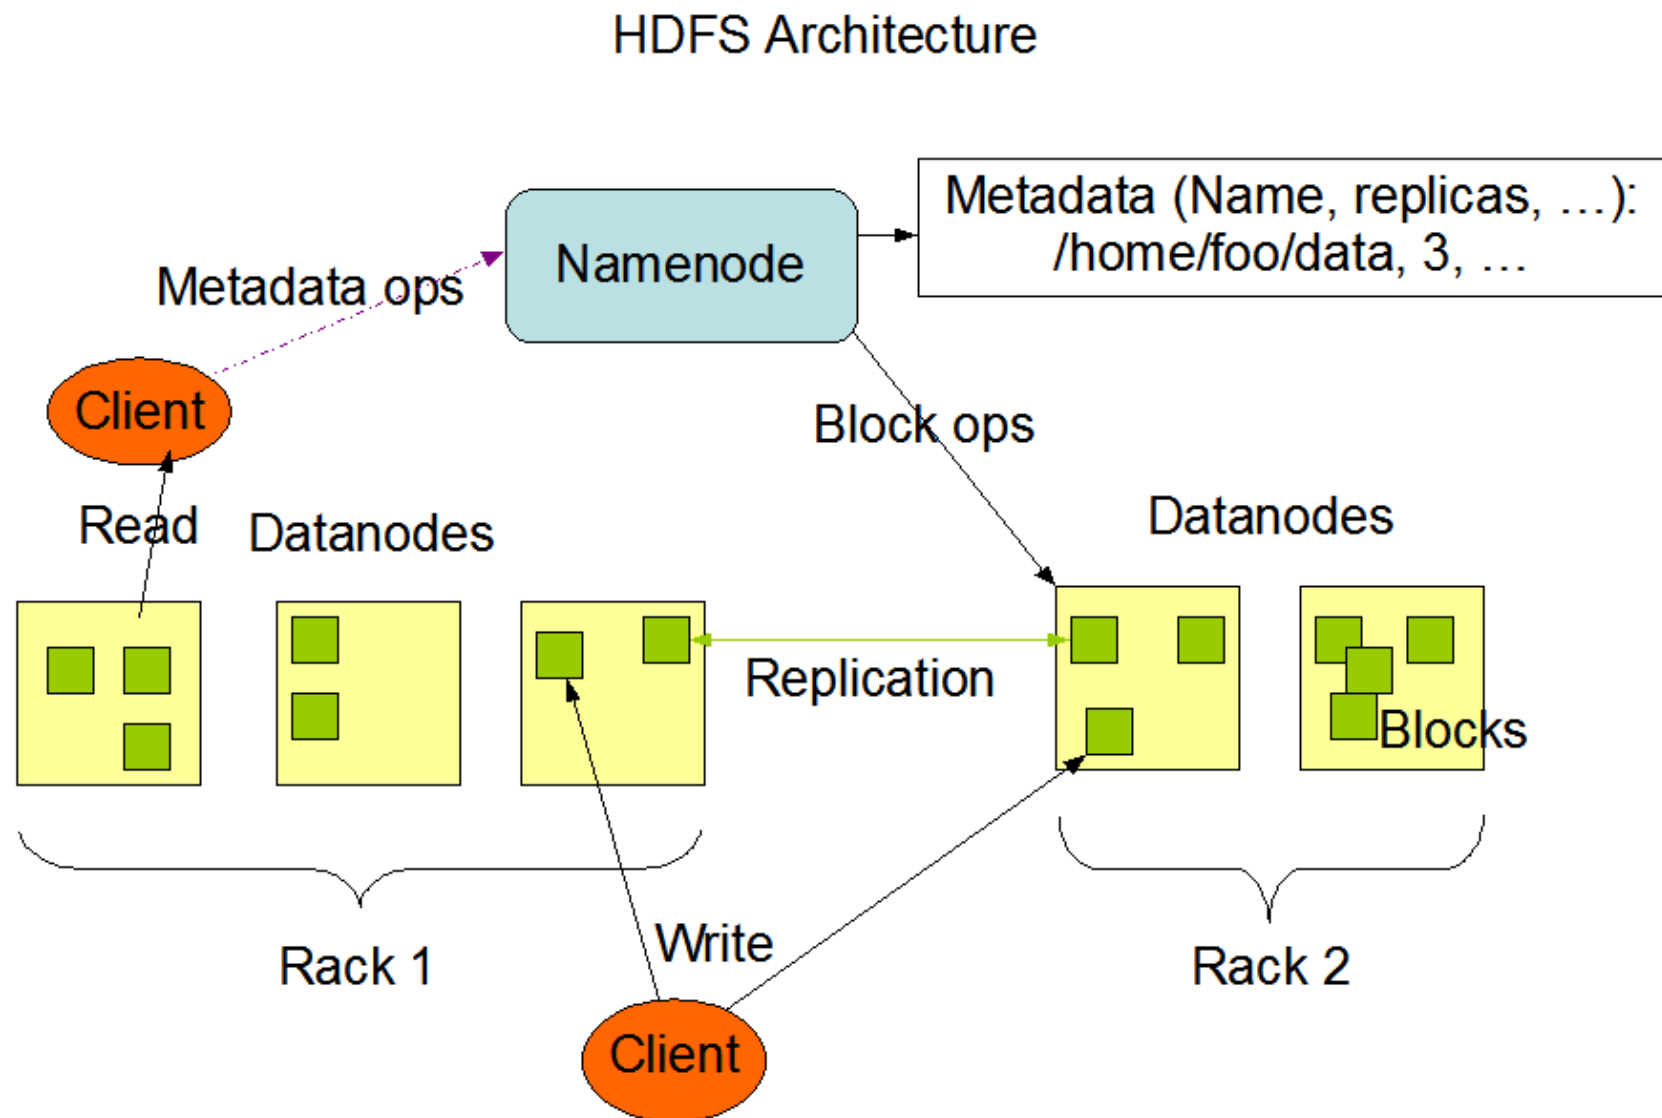
Agenda

- **HDFS**
- **Master/Slave architecture in HDFS**
- **The File System Namespace**
- **Data Replication**
- **Replica Selection**
- **Blockreport**
- **The Persistence of File System Metadata- EditLogs and FSImage**
- **Checkpoint**
- **Replication Pipelining**
- **Staging**
- **Space Reclamation**

HDFS

- HDFS is highly **fault-tolerant** and is designed to be deployed on low-cost hardware.
- HDFS provides **high throughput** access to application data and is suitable for applications that have large data sets.

Master/Slave architecture in HDFS



NameNode

- A master server that **manages the file system namespace** and **regulates access to files** by clients.
- The NameNode executes file system namespace operations like **opening, closing, and renaming files and directories**.
- It also determines the **mapping of blocks to DataNodes**.

DataNodes

- Manage storage attached to the nodes that they run on.
- Responsible for **serving read and write requests** from the file system's clients.
- Perform **block creation, deletion, and replication** upon instruction from the **NameNode**.

The File System Namespace

- The NameNode maintains the file system namespace. Any change to the file system namespace or its properties is recorded by the NameNode.
- Traditional hierarchical file organization.
- Exa:- **“/user/mydir/demo.csv”**
- A user or an application can create and remove files, move a file from one directory to another, or rename a file.
- HDFS does not yet implement user quotas.
- HDFS does not support hard links or soft links.
- Stores Replication Factor Information(number of replicas of a file that should be maintained by HDFS).

Data Replication

- HDFS stores each file as a **sequence of blocks**; all blocks in a file except the last block are the same size.
- The blocks of a file are replicated for fault tolerance.
- The block size and **replication factor are configurable per file**. An application can specify the number of replicas of a file.
- The replication factor **can be specified at file creation time** and can be changed later.

Data Replication


- Files in HDFS are write-once and have strictly one writer at any time.
- The **NameNode** makes all decisions regarding replication of blocks.
- It **periodically receives a Heartbeat** and a Blockreport from each of the DataNodes in the cluster.
- Receipt of a Heartbeat implies that the DataNode is functioning properly. A
- **Blockreport** contains a list of all blocks on a DataNode.

Data Replication

Block Replication

Namenode (Filename, numReplicas, block-ids, ...)
/users/sameerp/data/part-0, r:2, {1,3}, ...
/users/sameerp/data/part-1, r:3, {2,4,5}, ...

Datanodes



Replica Selection

- To minimize global bandwidth consumption and read latency, HDFS tries closest replica selection.
- If there exists a replica on the same rack as the reader node, then that replica is preferred to satisfy the read request.
- If an HDFS cluster spans multiple data centers, then a **replica that is resident in the local data center is preferred over any remote replica.**

Safemode

- On startup, the NameNode enters a special state called Safemode.
- Replication of data blocks does not occur when the NameNode is in the Safemode state.
- The NameNode receives Heartbeat and Blockreport messages from the DataNodes.
- Each block has a specified minimum number of replicas.

Safemode

- A block is considered **safely replicated** when the **minimum number of replicas** of that data block has **checked in with the NameNode**.
- After a configurable percentage of safely replicated data blocks checks in with the NameNode (plus an additional 30 seconds), the NameNode exits the Safemode state.
- It then determines the list of data blocks (if any) that still have **fewer than the specified number of replicas**.
- The NameNode then replicates these blocks to other DataNodes.

The Persistence of File System Metadata-EditLog

- The HDFS namespace is stored by the NameNode.
- The NameNode uses a **transaction log called the EditLog** to persistently **record every change** that occurs to file system metadata.
- For example, **creating a new file** in HDFS causes the NameNode to insert a record into the EditLog indicating this.
- Similarly, **changing the replication factor** of a file causes a new record to be inserted into the EditLog.
- EditLog are stored as a file in local host OS file system.

FsImage

- The **entire file system** namespace, including the **mapping of blocks to files and file system properties**, is stored in a file called the **FsImage**.
- The FsImage is stored as a file in the NameNode's local file system too.
- The NameNode **keeps an image** of the entire file system namespace and file Blockmap **in memory**.
- This key metadata item is designed to be compact, such that a NameNode with 4 GB of RAM is plenty to support a huge number of files and directories.

Checkpoint

- **When the NameNode starts up,**
 - it **reads** the FsImage and EditLog from disk,
 - **applies all the transactions** from the **EditLog** to the in-memory representation of the FsImage, and
 - **flushes out this new version** into a new FsImage on disk.
- **It can then truncate the old EditLog because its transactions have been applied to the persistent FsImage.**
- **This process is called a checkpoint.**
- **In the current implementation, a checkpoint only occurs **when the NameNode starts up**.**
- **Work is in progress to support periodic checkpointing in the near future.**

DataNode Storage

- The DataNode stores HDFS data in files in its local file system.
- The DataNode has no knowledge about HDFS files.
- **It stores each block of HDFS data in a separate file in its local file system.**
- The DataNode does not create all files in the same directory. Instead, it uses a heuristic to determine the optimal number of files per directory and creates subdirectories appropriately.
- It is not optimal to create all local files in the same directory because the local file system might not be able to efficiently support a huge number of files in a single directory.

Blockreport

- When a DataNode starts up, it scans through its local file system, generates **a list of all HDFS data blocks that correspond to each of these local files** and sends this report to the NameNode: this is the Blockreport.

The Communication Protocols in HDFS

- All HDFS communication protocols are layered on top of the TCP/IP protocol.
- A client establishes a **connection to a configurable TCP port** on the NameNode machine.
- It talks the **ClientProtocol with the NameNode**.
- The DataNodes talk to the NameNode using the **DataNode Protocol**.
- A Remote Procedure Call (RPC) abstraction wraps **both the Client Protocol and the DataNode Protocol**.
- By design, the NameNode never initiates any RPCs. Instead, it only responds to RPC requests issued by DataNodes or clients.

Staging- File Creation Operation in HDFS

- A client request to create a file does not reach the NameNode immediately.
- In fact, initially the HDFS client caches the file data into a temporary local file.
- Application writes are transparently redirected to this **temporary local file**.
- When the local file accumulates **data worth over one HDFS block size**, the client contacts the NameNode.
- The NameNode **inserts the file name** into the file system hierarchy and **allocates a data block** for it.

Staging- File Creation Operation in HDFS

- The NameNode responds to the client request with the **identity of the DataNode and the destination data block**.
- Then the **client flushes the block** of data from the local temporary file to the specified DataNode.
- When a file is closed, the remaining un-flushed data in the temporary local file is transferred to the DataNode.
- The client then tells the NameNode that the file is closed.
- At this point, the NameNode commits the file creation operation into a **persistent store**.
- If the NameNode dies before the file is closed, the file is lost.

Replication Pipelining- Writing data to HDFS from one DataNode to another DataNode

- When a client is writing data to an HDFS file, its **data is first written to a local file**.
- Suppose the HDFS file has a replication factor of three.
- When the local file **accumulates a full block** of user data, the **client retrieves a list of DataNodes from the NameNode**.
- This list contains the DataNodes that will host a replica of that block.
- The client then flushes the data block to the first DataNode.

Replication Pipelining- Writing data to HDFS from one DataNode to another DataNode

- The first DataNode starts receiving the data in small portions (4 KB), writes each portion to its local repository and transfers that portion to the second DataNode in the list.
- The second DataNode, in turn starts receiving each portion of the data block, writes that portion to its repository and then flushes that portion to the third DataNode.
- Finally, the third DataNode writes the data to its local repository.
- Thus, a **DataNode can be receiving data from the previous one in the pipeline and at the same time forwarding data to the next one in the pipeline.**
- Thus, the data is pipelined from one DataNode to the next.

Space Reclamation- What happens when you remove something?

- **There is two techniques to reclaim space in HDFS architecture:-**
 - File Deletes and Undeletes
 - Decrease Replication Factor

File Deletes

- When a file is deleted by a user or an application, it is not immediately removed from HDFS.
- Instead, **HDFS first renames it to a file in the /trash directory.**
- The file can be restored quickly as long as it remains in /trash.
- A file **remains in /trash** for a **configurable amount of time.**
- After the expiry of its life in /trash, the NameNode deletes the file from the HDFS namespace.
- The deletion of a file causes the **blocks associated with the file to be freed.**
- Note that there could be an appreciable time delay between the time a file is deleted by a user and the time of the corresponding increase in free space in HDFS.

File Undeletes

- A user can Undelete a file after deleting it as long as it remains in the **/trash directory**.
- If a user wants to undelete a file that he/she has deleted, he/she can navigate the /trash directory and retrieve the file.
- The /trash directory contains only the latest copy of the file that was deleted.
- The /trash directory is just like any other directory with one special feature: **HDFS applies specified policies to automatically delete files from this directory**.
- The current default policy is to **delete files from /trash that are more than 6 hours old**. In the future, this policy will be configurable through a well defined interface.

Decrease Replication Factor

- When the replication factor of a file is reduced, the **NameNode selects excess replicas that can be deleted.**
- The next **Heartbeat transfers this information to the DataNode.**
- The **DataNode then removes the corresponding blocks** and the corresponding free space appears in the cluster.
- There might be a time delay between the completion of the setReplication API call and the appearance of free space in the cluster.

Resources

- https://hadoop.apache.org/docs/current1/hdfs_design.html
- **Hadoop: The Definitive Guide**
 - Tom White (Author)
 - O'Reilly Media; 4th Edition.

